PRISMSPACE - INTELLIGENT MULTI-AGENT PLATFORM



Dept. of Computer Science & Engineering
College of Engineering, Karunagappally

Project Group 01

NOBIN SIJO	KNP23CS082
PRANAV P	KNP23CS085
PRASANTH P	KNP23CS086
SANKAR NARAYAN S	KNP23CS094

Project Coordinator & Guide
Dr. Leena Silvoster
Associate Professor
College of Engineering, Karunagappally

01. Introduction: PrismSpace Web Platform

PrismSpace is an interactive, **web-based multi-agent intelligence platform** and workspace website that allows users to orchestrate autonomous AI swarms directly from modern web browsers. It converts natural-language user goals into dependency-aware **Directed Acyclic Graphs (DAGs)** for parallel sub-task execution. Built as a responsive web application (Next.js 15, React 19) accessible from any browser with zero local setup, it features an interactive visual DAG swarm visualizer with real-time token and step streaming. The platform pairs Groq, OpenAI, and NVIDIA NIM with standardized tool connectivity via Model Context Protocol (MCP), while storing all session state locally in client-side IndexedDB (`PrismDB`) with human-in-the-loop approval gates for privacy and safety.

Problem Statement

Challenges in current AI-agent and workflow systems

02. Problem Statement: Core Challenges in AI Agent Systems

Current AI-agent systems suffer from monolithic model bottlenecks, rigid sequential execution, fragmented tool integrations, limited observability, hallucinated tool calls, and unsafe or privacy-sensitive actions. **PrismSpace** takes developer task prompts, workspace and session state, uploaded files, and MCP/HITL configurations as **inputs**; applies intent classification, supervised model/provider routing, DAG decomposition with topological scheduling, MCP-enabled ReAct execution, ORPO preference alignment, and human approval as its **method**; and produces **outputs** including verified artifacts, tool results, execution traces, quality and cost telemetry, and persistent session records.

03. Objectives

1. Decompose complex user goals into dependency-aware **DAGs** for structured and parallel task execution.
2. Orchestrate specialist agents and approved tools through **MCP** and bounded **ReAct** loops.
3. Align generated responses using **ORPO** preference optimization for quality, faithfulness, and safety.
4. Route tasks to suitable models and providers using supervised capability, cost, and latency signals.
5. Validate risky actions with **HITL** approval and preserve verified outputs with execution telemetry.

04. Literature Review

Research background, key findings, and identified gap

Multi-Objective Reward Generalization in Deep Reinforcement Learning

Citation	Summary	Advantages & Disadvantages
F. Cornalba, C. Disselkamp, D. Scassola, and C. Helf *Neural Computing and Applications* **Year:** 2024 **Publisher:** Springer Open Access	Investigates **Multi-Objective Deep Reinforcement Learning (MODRL)** for optimizing competing operational objectives (e.g., return vs. risk). Introduces reward generalization techniques that train pareto-optimal policies capable of dynamically adapting to shifting stakeholder preferences without retraining from scratch.	**Advantages:** ▶ Provides structured Pareto-optimal policy exploration across conflicting objectives. ▶ Generalizes across parameterized reward weights without full policy retraining. ▶ Directly informs multi-objective routing (latency vs. cost vs. quality) in PrismSpace. **Disadvantages:** ▶ High initial training sample complexity across multi-dimensional reward spaces. ▶ Evaluated primarily on continuous financial time-series data. ▶ Sensitive to hyperparameter tuning in policy entropy regularization.

Average Reward Adjusted Discounted Reinforcement Learning

Citation	Summary	Advantages & Disadvantages
M. Schneckenreither and G. Moser *Neural Computing and Applications* **Year:** 2025 **Publisher:** Springer Open Access	Proposes **Average-Reward Adjusted Discounted RL (Reinforcement Learning) (ARAD-RL)**, bridging the gap between discounted and average reward criteria in infinite-horizon **MDPs** (Markov Decision Processes). Addresses systematic value estimation bias in standard discounted Q-learning by dynamically subtracting the estimated gain rate during Bellman updates.	**Advantages:** ▶ Eliminates infinite-horizon return explosion while preserving policy convergence. ▶ Improves long-term steady-state throughput in continuous scheduling tasks. ▶ Theoretically sound formulation for continuous agent lifecycle optimization. **Disadvantages:** ▶ Requires accurate empirical estimation of average reward baseline. ▶ Slow convergence in highly non-stationary transition environments. ▶ Adds variance to early-stage **TD** (Temporal Difference) error gradient estimates.

Deep Reinforcement Learning for Multi-Agent Coordination

Citation	Summary	Advantages & Disadvantages
K. Aina and S. Ha *Artificial Life and Robotics* **Year:** 2026 **Publisher:** Springer Open Access	Develops decentralized **MARL** (Multi-Agent Reinforcement Learning) algorithms for multi-agent coordination in constrained environments. Employs **spatial-temporal observation encoding** and localized interaction graphs to coordinate joint movements and resolve deadlocks without requiring centralized global state knowledge.	**Advantages:** ▶ Resolves shared-resource contention and deadlocks via decentralized policies. ▶ Spatial-temporal graph attention scales linearly with agent density. ▶ Applicable to concurrent tool-use contention across worker agents. **Disadvantages:** ▶ Tested in physical robotics grids rather than software **API** (Application Programming Interface) environments. ▶ Vulnerable to communication packet dropouts and sensor latency. ▶ Does not account for semantic language-based coordination.

Open-Ended Coordination Using Modular Open Policies

Citation	Summary	Advantages & Disadvantages
D. Rother, J. Pajarinen, J. Peters, and T. H. Weisswange *Autonomous Agents and Multi-Agent Systems* **Year:** 2025 **Publisher:** Springer Open Access	Presents a framework for **open-ended ad hoc multi-agent coordination**. Utilizes **Modular Open Policies (MOP)** where dynamic agent teams dynamically infer unseen teammate skill representations and synthesize blended coordination strategies without requiring pre-coordinated joint training.	**Advantages:** ▶ Enables zero-shot ad hoc coordination with unfamiliar teammate agent policies. ▶ Modular policy decomposition allows plug-and-play agent role assignment. ▶ Directly adopted in PrismSpace for dynamic worker swarm scaling. **Disadvantages:** ▶ Requires comprehensive pre-trained skill policy libraries for base roles. ▶ Evaluated primarily on benchmark games (Overcooked, **LBF** (Level-Based Foraging)). ▶ Action blending may introduce transient sub-optimal execution steps.

Improving Scalability of **MADRL** (Multi-Agent Deep Reinforcement Learning) with Suboptimal Human Knowledge

Citation	Summary	Advantages & Disadvantages
D. Liu, F. Ren, J. Yan, G. Su, W. Gu, and S. Kato *Autonomous Agents and Multi-Agent Systems* **Year:** 2026 **Publisher:** Springer Open Access	Addresses sample inefficiency in large-scale **Multi-Agent DRL** (Deep Reinforcement Learning) by integrating heuristic human guidance into exploration. Proposes an adaptive filtering mechanism that extracts actionable macro-actions from noisy or sub-optimal human demonstration traces, accelerating multi-agent policy learning.	**Advantages:** ▶ Accelerates policy exploration by up to 4x using noisy demonstration priors. ▶ Robust to suboptimal or conflicting human heuristic rules. ▶ Informs PrismSpace's bootstrap initialization from logged developer sessions. **Disadvantages:** ▶ Quality of learned policies depends on initial heuristic coverage. ▶ Requires complex threshold tuning to transition away from demonstrations. ▶ High memory footprint for maintaining dual replay buffers.

ε-Retraining Reinforcement Learning Algorithms

Citation	Summary	Advantages & Disadvantages
L. Marzari, C. Liu, P. L. Donti, and E. Marchesini *Autonomous Agents and Multi-Agent Systems* **Year:** 2026 **Publisher:** Springer Open Access	Investigates bounded policy fine-tuning through ε-**retraining**, preventing catastrophic forgetting when deep **RL** agents adapt to non-stationary environments. Establishes formal Lipschitz bounds on policy drift, ensuring safe continual learning in deployed agent pipelines.	**Advantages:** ▶ Provides mathematical guarantees against catastrophic policy divergence. ▶ Enables continuous telemetry-driven model fine-tuning in production. ▶ Crucial for stability in PrismSpace's ongoing user interaction learning. **Disadvantages:** ▶ Constrained optimization increases computation per gradient step. ▶ Conservative update bounds may slow adaptation to drastic distribution shifts. ▶ Theoretical guarantees rely on local smoothness assumptions.

Deep Multiagent Reinforcement Learning: Challenges and Directions

Citation	Summary	Advantages & Disadvantages
A. Wong, T. Bäck, A. V. Kononova, and A. Plaat *Artificial Intelligence Review* **Year:** 2023 **Publisher:** Springer Open Access	Surveys **Deep Multiagent Reinforcement Learning (MARL)**, formalizing foundational challenges including non-stationarity, multi-agent credit assignment, and combinatorial action explosion. Evaluates **Centralized Training with Decentralized Execution (CTDE)** and communication architectures for modular orchestration.	**Advantages:** ▶ Comprehensive taxonomy of credit assignment and **CTDE** paradigms. ▶ Establishes formal theoretical criteria for communication protocols. ▶ Provides the theoretical foundation for PrismSpace's modular agent architecture. **Disadvantages:** ▶ Survey paper without empirical experimental evaluation. ▶ Focuses on discrete game simulations rather than **API**-driven software tasks. ▶ Does not cover **LLM** reasoning agents or human-in-the-loop approval gates.

HIAT (Human-in-the-Loop with Auxiliary Tasks): Human-in-the-Loop **RL** with Auxiliary Tasks

Citation	Summary	Advantages & Disadvantages
B. Niu, B. Luo, Y. Cui, X. Xu, Y. Zhao, and Y. Feng *Artificial Intelligence Review* **Year:** 2025 **Publisher:** Springer Open Access	Introduces **HIAT (Human-in-the-Loop with Auxiliary Tasks)**, a framework leveraging human intervention signals as auxiliary self-supervised learning objectives. Models human corrections as feedback gradients that shape internal feature representations, minimizing human intervention frequency over time.	**Advantages:** ▶ Maximizes learning efficiency from scarce human intervention signals. ▶ Significantly reduces human cognitive load through auxiliary reward shaping. ▶ Informs PrismSpace's approval gate telemetry and learning from user edits. **Disadvantages:** ▶ Requires consistent human feedback to avoid biased auxiliary representations. ▶ Intervention timing latency can cause temporal mis-alignment in **RL** updates. ▶ Higher computational overhead during auxiliary task representation training.

ORPO (Odds Ratio Preference Optimization): Monolithic Preference Optimization without Reference Model

Citation	Summary	Advantages & Disadvantages
J. Hong, N. Lee, and J. Thorne *EMNLP* / *arXiv:2403.07691* **Year:** 2024 **Area: LLM** Preference Alignment	**Core Idea (Single-Phase Alignment):** Instead of multi-step **RLHF**, **ORPO** teaches desired answers while penalizing bad answers in one unified equation: $$\mathcal{L}_{\text{ORPO}} = \mathcal{L}_{\text{SFT}} + \lambda \cdot \mathcal{L}_{\text{OR}}$$ $$\mathcal{L}_{\text{OR}} = -\log \sigma \left(\log \frac{\text{odds}(y_w \mid x)}{\text{odds}(y_l \mid x)} \right)$$ **Key Terms:** ▶ $\mathcal{L}_{\text{SFT}}$: Base supervised fine-tuning loss (correct code/syntax). ▶ $\mathcal{L}_{\text{OR}}, \lambda$: Odds ratio penalty term and penalty weight. ▶ x, y_w, y_l: Prompt, winning action path, and losing action path.	**Advantages:** ▶ Reference-model-free: eliminates maintaining frozen π_{ref} in memory. ▶ Monolithic single-phase training avoids unstable **PPO** policy optimization. ▶ Directly aligns PrismSpace agent trajectories on paired feedback (*chosen* vs. *rejected*). **Disadvantages:** ▶ Requires high-quality pairwise preference datasets (y_w, y_l). ▶ Sensitive to λ hyperparameter balance.

Reinforcement Online Learning to Rank with Unbiased Reward Shaping

Citation	Summary	Advantages & Disadvantages
S. Zhuang, Z. Qiao, and G. Zuccon *Information Retrieval Journal* **Year:** 2023 **Publisher:** Springer Open Access	Develops an **Online Learning to Rank (OLTR)** framework using counterfactual reinforcement learning. Employs **inverse propensity weighting (IPW)** and unbiased reward shaping to correct for position, presentation, and click bias in user interaction logs during continuous ranker optimization.	**Advantages:** ▶ Guarantees unbiased gradient estimation from observational implicit feedback. ▶ Operates entirely online without requiring manual document relevance labels. ▶ Directly informs PrismSpace's context retrieval ranking feedback loop. **Disadvantages:** ▶ High propensity score variance can destabilize gradient updates. ▶ Requires large volumes of query-click logs to converge. ▶ Assumes position-based user click models that may not hold in code search.

Knowledge Learning for Securing Workflow Scheduling in **MEC** (Mobile Edge Computing)

Citation	Summary	Advantages & Disadvantages
T. M. Ghazal, A. E. Awouda, M. K. Hasan, A. H. A. Rahman, S. Islam, R. A. Saeed, H. Elshafie, A. Iqbal, and A. H. Hussein *Mobile Networks and Applications* **Year:** 2025 **Publisher:** Springer Open Access	Proposes a **Knowledge-Guided Reinforcement Learning (KG-RL)** scheduler for complex workflow **DAGs** in distributed edge-cloud environments. Integrates security constraint validation directly into the state transition model, preventing insecure task placement while optimizing makespan.	**Advantages:** ▶ Enforces strict security boundary constraints during automated task scheduling. ▶ Effectively handles multi-dependency Directed Acyclic Graph (**DAG**) workflows. ▶ Informs PrismSpace's sandboxed tool execution and dependency resolution. **Disadvantages:** ▶ Constraint checking introduces scheduling latency for massive **DAGs**. ▶ Assumes static security classification matrices across network nodes. ▶ Limited evaluation on non-deterministic tool failure scenarios.

Energy-Aware Cooperative Multi-Fitness Evolutionary Workflow Scheduling

Citation	Summary	Advantages & Disadvantages
P. Barredo and J. Puente *Natural Computing* **Year:** 2025 **Publisher:** Springer Open Access	Presents a **Cooperative Multi-Fitness Evolutionary Algorithm (CMFEA)** for multi-objective workflow scheduling. Optimizes trade-offs between execution time, monetary compute cost, and energy consumption using island-model cooperative coevolution across heterogeneous worker pools.	**Advantages:** ▶ Pareto-optimal trade-offs across makespan, financial cost, and energy. ▶ Cooperative sub-population decomposition prevents premature local convergence. ▶ Provides heuristics for PrismSpace's multi-agent model provider selection. **Disadvantages:** ▶ Evolutionary search has higher runtime latency than single-pass ML heuristics. ▶ Cannot adapt instantaneously to real-time **API** rate limit drops. ▶ Requires parameter tuning across mutation, crossover, and migration rates.

Optimized Task Scheduling Using Hybrid **DL** (Deep Learning) and Metaheuristics

Citation	Summary	Advantages & Disadvantages
A. Heirati, M. Fartash, and M. Khalily-Dermany *Neural Processing Letters* **Year:** 2026 **Publisher:** Springer Open Access	Combines **Deep Learning (LSTM (Long Short-Term Memory))** workload prediction with metaheuristic optimization (Harris Hawks Optimization) for proactive workflow dispatching. Predicts task execution durations and memory bottlenecks in advance to prevent node overloading and minimize queue waiting times.	**Advantages:** ▶ Proactive workload forecasting prevents sudden compute node thrashing. ▶ Hybrid pipeline balances search exploration with predictive dispatching. ▶ Relevant for predicting token latency and cost before agent invocation. **Disadvantages:** ▶ Two-stage pipeline introduces inference overhead before scheduling decisions. ▶ **LSTM** prediction accuracy degrades under sudden irregular workload spikes. ▶ Complex offline model retraining requirements.

Cost-Aware Workflow Offloading in Edge-Cloud Computing via **GA** (Genetic Algorithm)

Citation	Summary	Advantages & Disadvantages
S. Abdi, M. Ashjaei, and S. Mubeen *The Journal of Supercomputing* **Year:** 2024 **Publisher:** Springer Open Access	Proposes a **Genetic Algorithm-based workflow offloading** framework considering heterogeneous network bandwidth, deadline constraints, and monetary costs. Implements chromosome repair operators to preserve **DAG** precedence validity throughout mutation and crossover.	**Advantages:** ▶ Guarantees valid **DAG** dependency structures via specialized repair operators. ▶ Explicitly accounts for network transfer latency between pipeline stages. ▶ Informs subtask payload transfer between worker agents in PrismSpace. **Disadvantages:** ▶ Computationally heavy for interactive real-time user query response. ▶ Assumes deterministic network transfer bandwidth in cloud environments. ▶ Does not model dynamic agent failures or retries.

Dynamic Role Discovery and Assignment in Multi-Agent Task Decomposition

Citation	Summary	Advantages & Disadvantages
Y. Xia, J. Zhu, and L. Zhu *Complex & Intelligent Systems* **Year:** 2023 **Publisher:** Springer Open Access	Proposes **dynamic role discovery and assignment** that decomposes complex high-level objectives into lower-level subtasks. Uses learned persona embeddings and mutual information to adaptively allocate subtasks according to situational agent capabilities, eliminating rigid static assignments.	**Advantages:** ▶ Enables flexible subtask allocation without hardcoded role hierarchies. ▶ Hierarchical decomposition prevents single-agent cognitive context overload. ▶ Serves as core theoretical basis for PrismSpace's dynamic agent routing. **Disadvantages:** ▶ Role discovery clustering introduces upfront computational overhead. ▶ Performance depends on representation vector accuracy for subtasks. ▶ Evaluated primarily on synthetic collaborative coordination benchmarks.

Decision Variable Contribution Based Adaptive Multi-Objective Scheduling

Citation	Summary	Advantages & Disadvantages
J. Li, L. Xing, W. Zhong, Z. Cai, and F. Hou *Complex & Intelligent Systems* **Year:** 2023 **Publisher:** Springer Open Access	Introduces a **Decision Variable Contribution (DVC)** adaptive mechanism for evolutionary multi-objective workflow scheduling. Evaluates the differential sensitivity of individual decision variables to focus evolutionary search operators on high-impact scheduling bottlenecks.	**Advantages:** ▶ Accelerates Pareto-front convergence by prioritizing high-variance variables. ▶ Dynamically adapts search intensity across workflow bottleneck stages. ▶ Useful for pinpointing high-latency critical path tools in agent workflows. **Disadvantages:** ▶ Sensitivity calculation adds per-generation matrix computation cost. ▶ Can over-concentrate on dominant objectives at the expense of diversity. ▶ Requires calibrated normalization across widely different metric units.

GHQ: Grouped Hybrid Q-Learning for Heterogeneous Multi-Agent RL

Citation	Summary	Advantages & Disadvantages
X. Yu, Y. Lin, X. Wang, S. Han, and K. Lv *Complex & Intelligent Systems* **Year:** 2024 **Publisher:** Springer Open Access	Develops **Grouped Hybrid Q-learning (GHQ)** for cooperative heterogeneous multi-agent systems. Clusters agents with similar capabilities into functional cohorts that share value approximations while maintaining decoupled action execution heads for role specialization.	**Advantages:** ▶ Scales value estimation across large heterogeneous swarms via cohort grouping. ▶ Reduces joint action space dimensionality from exponential to block-linear. ▶ Directly applicable to categorizing **LLM** worker agent skill specializations. **Disadvantages:** ▶ Cohort clustering boundaries can be rigid if agent capabilities evolve. ▶ Intra-group credit assignment remains non-trivial in dense interactions. ▶ Requires domain knowledge to establish initial grouping heuristics.

Decentralized Multi-Agent Cooperation via Adaptive Partner Modeling

Citation	Summary	Advantages & Disadvantages
C. Xu, J. Wang, X. Zhu, Y. Yue, W. Zhou, Z. Liang, and D. Wojtczak *Complex & Intelligent Systems* **Year:** 2024 **Publisher:** Springer Open Access	Tackles environmental non-stationarity in decentralized multi-agent cooperation using **Adaptive Partner Modeling (APM)**. Instead of relying on expensive continuous inter-agent broadcast communication, agents infer partner policies and intended actions from local observation histories.	**Advantages:** ▶ Minimizes inter-agent communication bandwidth to near zero during execution. ▶ Stabilizes decentralized joint execution despite teammate policy shifts. ▶ Directly adopted in PrismSpace for zero-communication agent coordination. **Disadvantages:** ▶ Partner inference error propagates if historical observations are noisy. ▶ Inference networks introduce local computational and memory overhead. ▶ Assumes partners follow relatively stationary parameterized policy spaces.

A Responsible AI Framework: Pipeline Contextualisation

Citation	Summary	Advantages & Disadvantages
E. Vyhmeister, G. G. Castañé, P.-O. Östberg, and S. Thevenin *AI and Ethics* **Year:** 2023 **Publisher:** Springer Open Access	Formulates an operational **Responsible AI Framework** providing pipeline contextualisation across the full model lifecycle. Specifies auditable checkpoints, data lineage tracking, and bias mitigation protocols for autonomous AI decision-making systems.	**Advantages:** ▶ Holistic governance framework linking technical architecture to ethical safety. ▶ Defines concrete telemetry points for data provenance and action auditing. ▶ Guides PrismSpace's execution audit logging and human inspection UI. **Disadvantages:** ▶ Qualitative framework that requires custom translation into concrete code gates. ▶ Audit logging introduces storage overhead and slight execution latency. ▶ Does not provide automated formal verification of **LLM** reasoning traces.

Meaningful Human Control: Actionable Properties for AI Development

Citation	Summary	Advantages & Disadvantages
L. Cavalcante Siebert, M. L. Lupetti, E. Aizenberg, N. Beckers, A. Zgonnikov, H. Veluwenkamp, D. Abbink, E. Giaccardi, G.-J. Houben, C. M. Jonker, J. van den Hoven, D. Forster, and R. L. Lagendijk *AI and Ethics* **Year:** 2023 **Publisher:** Springer Open Access	Formalizes the philosophical concept of **Meaningful Human Control (MHC)** into actionable software design properties: **Tracking** (system responsive to human moral reasons) and **Tracing** (system actions causally attributable to accountable human supervisors).	**Advantages:** ▶ Translates abstract governance principles into rigorous software specifications. ▶ Establishes tracking and tracing requirements for autonomous tool execution. ▶ Directly guides the design of PrismSpace's Human-in-the-Loop (**HITL**) approval gates. **Disadvantages:** ▶ High human intervention frequency can degrade automated agent throughput. ▶ Defining actionable reason boundaries requires careful domain calibration. ▶ Assumes human supervisor maintains full comprehension of complex agent plans.

CPN-HRL (Cloud-Private-Network Hierarchical Reinforcement Learning): Hierarchical **DRL** for Priority-Aware Task Scheduling

Citation	Summary	Advantages & Disadvantages
Q. Yue, L. Tian, X. Feng, Z. Yuan, S. Wei, J. Xia, B. Chen, X. Song, and Y. Yao *Journal of King Saud University CIS* **Year:** 2026 **Publisher:** Springer Open Access	Presents **CPN-HRL**, a Hierarchical Reinforcement Learning framework for priority-aware task scheduling in Cloud-Edge-End networks. Uses a macro-policy to partition task **DAGs** by priority and latency budgets, while micro-policies schedule subtasks onto optimal compute nodes.	**Advantages:** ▶ Hierarchical policy decomposition effectively handles complex task **DAGs**. ▶ Guarantees priority-aware quality of service under tight latency deadlines. ▶ Serves as primary reference for PrismSpace's **DAG** workflow scheduling engine. **Disadvantages:** ▶ Multi-tier policy architecture requires synchronous training coordination. ▶ Edge node failure requires immediate micro-policy replanning. ▶ Evaluated in simulated telecom cloud-edge topologies.

Automatic Discovery and Description of Human Planning Strategies

Citation	Summary	Advantages & Disadvantages
J. Skirzyński, Y. Jain, and F. Lieder *Behavior Research Methods* **Year:** 2023 **Publisher:** Springer Open Access	Develops an algorithmic framework for **automatically discovering and formalizing planning strategies** from human behavior traces. Models task planning as resource-rational search over Directed Graphical Models, producing interpretable decision trees that characterize human heuristic planning.	**Advantages:** ▶ Extracts interpretable, human-like heuristic planning strategies automatically. ▶ Quantifies trade-offs between cognitive search effort and decision quality. ▶ Informs PrismSpace's hierarchical prompt decomposition and agent planning. **Disadvantages:** ▶ Search space for strategy discovery grows exponentially with problem depth. ▶ Human behavioral experiments conducted in stylized cognitive testbeds. ▶ Static decision trees struggle to capture fluid **LLM** reasoning trajectories.

Data-Intensive Workflow Scheduling Based on **DRL** in Multi-Clouds

Citation	Summary	Advantages & Disadvantages
S. Zhang, Z. Zhao, C. Liu, and S. Qin *Journal of Cloud Computing* **Year:** 2023 **Publisher:** Springer Open Access	Proposes a **Deep Q-Network (DQN)** scheduler optimized for data-intensive workflows across multi-cloud environments. Explicitly models dataset placement and cross-datacenter transfer latency to minimize both workflow completion time and data egress billing.	**Advantages:** ▶ Optimizes data locality to prevent massive cross-cloud egress overhead. ▶ Adaptive reward function balances monetary cost and makespan constraints. ▶ Directly relevant to caching and embedding retrieval across distributed agents. **Disadvantages:** ▶ Assumes deterministic file sizes and network bandwidth availability. ▶ Q-network value overestimation requires double-**DQN** correction. ▶ State representation size scales with the total number of cloud datacenters.

RL-Based Task Scheduling for Sustainable Federated Cloud Computing

Citation	Summary	Advantages & Disadvantages
Z. Wang, S. Chen, L. Bai, J. Gao, J. Tao, R. R. Bond, and M. D. Mulvenna *Journal of Cloud Computing* **Year:** 2023 **Publisher:** Springer Open Access	Develops an **Actor-Critic RL scheduling framework** for federated cloud datacenters that incorporates renewable energy availability into scheduling policies. Dynamically routes compute tasks to datacenters with green energy surpluses without violating user SLA deadlines.	**Advantages:** ▶ Incorporates environmental sustainability and carbon metrics into scheduling. ▶ Actor-Critic policy gradient handles continuous action spaces effectively. ▶ Provides conceptual foundation for multi-attribute provider routing. **Disadvantages:** ▶ Requires real-time telemetry from green energy generation sources. ▶ Trade-off between carbon footprint and network transmission latency. ▶ Sensitive to reward scaling factors between SLA penalties and green rewards.

Hybrid Dolphin Swarm Sparrow Search Multi-Objective Workflow Scheduling

Citation	Summary	Advantages & Disadvantages
S. Subashree, M. Rajakumaran, G. Pushpa, and S. Senthilkumar *Journal of Cloud Computing* **Year:** 2025 **Publisher:** Springer Open Access	Proposes a hybrid bio-inspired metaheuristic combining **Dolphin Swarm Optimization (DSA)** and **Sparrow Search Algorithm (SSA)** for multi-objective cloud workflow scheduling. Optimizes execution makespan, virtual machine utilization, and monetary execution expenditure.	**Advantages:** ▶ Combines global exploration of DSA with local exploitation of SSA. ▶ Demonstrates superior Pareto hypervolume compared to standard **GA** and **PSO** (Particle Swarm Optimization). ▶ Provides strong comparative baselines for workflow optimization algorithms. **Disadvantages:** ▶ Metaheuristic iteration runtime is too slow for sub-second reactive routing. ▶ Multiple tuning parameters (predation rate, vocalization index) required. ▶ Lacks theoretical convergence rate bounds for stochastic optimization.

Research Gap

What the surveyed methods already do

MARL works address **CTDE**, role grouping, and ad-hoc coordination. **ORPO** aligns **LLM** preferences without reference models, while **HIAT** leverages human interventions as auxiliary **RL** signals. **CPN-HRL** and **DAG** schedulers optimize cloud task allocation, and **MHC** formalizes tracking and tracing properties.

Limitations observed in those papers

Coordination benchmarks focus on gaming (**SMAC**, Overcooked) rather than filesystem/**API** tooling. **ORPO** lacks grounding in **MCP ReAct** tool traces, while **HIAT** targets continuous robotic control instead of code execution. Cloud schedulers assume static **DAGs** or slow metaheuristics ill-suited for real-time **LLM** routing.

Why PrismSpace is required

No reviewed system jointly grounds tool observations, aligns policies on agent trajectories, dynamically routes heterogeneous **LLMs**, and enforces human-in-the-loop safety gates within a unified developer workspace.

03. Comparative Analysis: Existing Platforms vs. PrismSpace

Simple comparison: Comparing PrismSpace with conversational AI, autonomous-agent, and multi-agent platforms.

Capability Dimension	ChatGPT Web (OpenAI)	AutoGPT (Autonomous Loop)	CrewAI (Multi-Agent Swarm)	PrismSpace (Ours) Local Multi-Agent Platform
Execution Topology	One conversation at a time	Repeats steps automatically; may not stop	Agents follow a fixed sequence	**DAG (Directed Acyclic Graph) Scheduling:** Runs dependent tasks in parallel, with up to three workers
Model Strategy & Routing	Mostly one provider and model	Usually one model with custom scripts	Model selected for each agent role	**Multi-Model Router:** Selects a suitable provider using task needs, speed, and cost
Tool Connectivity Protocol	Built-in plugins and function calls	Custom scripts and shell commands	Tool wrappers with fixed formats	**MCP (Model Context Protocol):** A common interface for tools such as Gmail, Drive, files, and web services
Observability & Streaming	Shows answer text; internal steps are hidden	Terminal logs only	Text traces from agent execution	**Live Run Status + Workflow Graph:** Shows task progress and each node's state
Safety & Quality Control	Provider filters; no user approval step	Fixed limits; limited output checking	Agent reviews its own output	**HITL (Human-in-the-Loop) + ORPO (Odds Ratio Preference Optimization):** Pauses risky tasks and ranks preferred outputs
Persistence & Privacy	Data stored in the provider's cloud	Local files without a structured database	Cloud or temporary in-memory state	**Local-First PrismDB:** Stores sessions and files in the browser's IndexedDB database

PrismSpace summary: It combines dependency-aware task planning, shared tool connections, provider selection, visible progress, and local browser storage in one workspace.

Algorithm

Step-by-step execution of the PrismSpace multi-agent workflow

Algorithm: Step 1 – Understand the User Prompt

Step 1: Prompt Parsing and Structured Representation

1.1 Read the user's natural-language prompt and identify the main goal.

1.2 Retrieve persistent session context and workspace history from browser IndexedDB (`PrismDB`):

- **1.2.1** Load relevant conversation history, saved preferences, and agent state.
- **1.2.2** Resolve referenced files, repositories, and available permissions.

1.3 Build a structured task record containing the goal, constraints, context, allowed tools, and expected output.

- **1.3.1** Record deadlines, output format, and other task constraints.
- **1.3.2** Attach the MCP tools and permissions required for execution.

1.4 Perform input boundary and security screening:

- **1.4.1** Check input size and screen for prompt-injection patterns.
- **1.4.2** Mark file, email, or other state-changing actions for user approval.

Step 2: Intent Classification and Capability Check

2.1 Classify the prompt into a functional domain such as automation, document work, research, or coding:

- **2.1.1** Extract semantic and keyword features from the prompt.
- **2.1.2** Use the intent classifier to select the most likely domain.

2.2 Profile required system capabilities and operational risk levels:

- **2.2.1** Identify required resources, such as web search, file access, or terminal execution.
- **2.2.2** Separate read-only work from actions that need approval.

2.3 Select optimal execution pathway:

- **2.3.1** Send simple requests directly to one model.
- **2.3.2** Send complex requests to the DAG planner for decomposition.

Step 3: DAG Decomposition and Model Routing

3.1 Deconstruct objective into Directed Acyclic Graph $\mathcal{G} = (V, E)$:

- **3.1.1** Create one vertex for each atomic task and define its expected output.
- **3.1.2** Connect tasks with directed edges when one task depends on another's output.
- **3.1.3** Hold a dependent task until all required inputs are available.

3.2 Compute topological task execution schedule $\pi = \mathrm{TopologicalSort}(\mathcal{G})$:

- **3.2.1** Find independent tasks that can run in parallel ($K \leq 3$).
- **3.2.2** Release dependent tasks when their parent tasks finish.

3.3 Select a model for each sub-task by balancing cost, latency, and quality:

$$\text{Provider}^* = \arg\min_m \left(\lambda_{\text{cost}} \, \text{Cost}_m + \lambda_{\text{latency}} \, \text{Latency}_m \right) \quad \text{s.t. } \text{Quality}_m \geq \tau$$

Meaning: Choose the lowest-cost and lowest-latency model that still meets the required quality threshold.

- **3.3.1** Use fast backends for simple or latency-sensitive tasks.
- **3.3.2** Use stronger models for reasoning, research, and code generation.

Step 4: Role-Based Task Assignment and Dispatching

4.1 Check the specialist-agent registry and their capability profiles:

- **4.1.1** Review each agent's role, instructions, and assigned tools.
- **4.1.2** Check worker availability before dispatch.

4.2 Match each sub-task to the agent with the closest capability profile:

$$\text{Agent}^{*} = \arg\max_{a \in \mathcal{A}} \ \text{CosineSimilarity}\big(\mathbf{Req}(v_i),\, \mathbf{Cap}(a)\big)$$

Meaning: Compare the task requirements with each agent's capabilities and choose the best match.

- **4.2.1** Send information-gathering tasks to the Researcher.
- **4.2.2** Send software tasks to the Coder.
- **4.2.3** Send checking and policy tasks to the Validator.

4.3 Manage task lifecycle states in dispatcher queue:

- **4.3.1** Keep tasks in `PENDING` until their dependencies finish.
- **4.3.2** Move ready tasks to `READY` and dispatch them.

Step 5: Agent Decision Cycle and Inter-Agent Communication

5.1 Run a ReAct cycle for each dispatched agent:

$$\text{State } s_t \longrightarrow \textbf{Observe} \longrightarrow \textbf{Reason} \longrightarrow \textbf{Act} \longrightarrow \textbf{Reflect}$$

Meaning: Observe the task, reason about the next step, act with a tool or response, and reflect on the result.

- **5.1.1** **Observe:** Read the task, constraints, and parent outputs.
- **5.1.2** **Reason:** Choose the next operation.
- **5.1.3** **Act:** Call an MCP tool or generate an answer.
- **5.1.4** **Reflect:** Check the result and correct local errors.

5.2 Transmit intermediate artifacts via DAG edge message passing:

- **5.2.1** Pass verified outputs from completed tasks to their dependent tasks.
- **5.2.2** Pass only relevant artifacts to limit context growth.

5.3 Synchronize shared state via local blackboard memory:

- **5.3.1** Save task state, conversation turns, and file references in `PrismDB`.
- **5.3.2** Notify dependent workers when a task completes.

Step 6: Specialist Execution Workflows – Planner & Assistant

6.1 Execute Planner Agent workflow (System Architect & Coordinator):

- **6.1.1** Read the objective and relevant context from `PrismDB`.
- **6.1.2** Split the objective into independent and dependent sub-goals.
- **6.1.3** Build the DAG and add capability tags.
- **6.1.4** Select models and enqueue the topological schedule.

6.2 Execute Assistant Agent workflow (Productivity & Communication):

- **6.2.1** Retrieve preferences and reference documents.
- **6.2.2** Create a structured draft or summary.
- **6.2.3** Use Google Workspace tools when required.
- **6.2.4** Verify the draft and request approval before sending it.

Step 7: Specialist Execution Workflows – Researcher, Coder & Validator

7.1 Execute Researcher Agent workflow (Information Retrieval & Synthesis):
- **7.1.1** Create search queries and use the Browser MCP server.
- **7.1.2** Extract relevant information from retrieved documents.
- **7.1.3** Remove duplicates and prepare a cited report.

7.2 Execute Coder Agent workflow (Software Synthesis & Execution):
- **7.2.1** Inspect the repository and edit the required source files.
- **7.2.2** Run tests or scripts in the terminal and inspect errors.
- **7.2.3** Fix syntax and runtime problems using test feedback.

7.3 Execute Validator Agent workflow (Quality Auditing & Safety Gate):
- **7.3.1** Check outputs against safety, format, and functional requirements.
- **7.3.2** Return **PASS** for valid output or **FAIL** with diagnostics.

Step 8: Standardized Tool Execution via Model Context Protocol

8.1 Interface agents to external tools via Model Context Protocol (MCP):

- **8.1.1** Use JSON-RPC for consistent client-server communication.
- **8.1.2** Discover tool schemas, inputs, and return types at runtime.
- **8.1.3** Keep agent logic independent of tool implementation details.

8.2 Orchestrate multi-domain tool execution ecosystem:

- **8.2.1** **System & Dev Tools:** Filesystem, terminal, and GitHub.
- **8.2.2** **Productivity Tools:** Gmail and Google Drive.
- **8.2.3** **Web Intelligence:** Browser search and document retrieval.

8.3 Enforce Human-in-the-Loop (HITL) authorization policy:

- **8.3.1** Allow read-only queries to run automatically.
- **8.3.2** Request approval for file changes, email sending, or other state changes.

Step 9: Context Compaction and Shared Memory Management

9.1 Address context window constraints and token consumption limits:

- **9.1.1** Keep the active context within the model's token limit.
- **9.1.2** Reduce unnecessary history to control latency and attention cost.

9.2 Compact the context by keeping verified outputs and removing stale or repeated content.

- **9.2.1** Forward only verified parent outputs to dependent tasks.
- **9.2.2** Remove redundant logs, drafts, and intermediate reasoning.

9.3 Persist global execution state using client-side database (`PrismDB`):

- **9.3.1** Save session turns, task graphs, and tool logs in IndexedDB.
- **9.3.2** Preserve local state for privacy and recovery after failure.

Step 10: Single-Stage Preference Alignment (ORPO) and Verification

10.1 Train the response-quality model with chosen and rejected examples using ORPO:

$$\mathcal{L}_{\text{ORPO}} = \mathcal{L}_{\text{SFT}} + \lambda_{\text{OR}} \cdot \mathcal{L}_{\text{OR}} \quad \text{where} \quad \mathcal{L}_{\text{OR}} = -\log \sigma \left(\log \left[\frac{\text{odds}_\theta(y_w \mid x)}{\text{odds}_\theta(y_l \mid x)} \right] \right)$$

Meaning: Combine normal supervised loss with a preference loss that favors the chosen response y_w over the rejected response y_l.

10.1.1 Increase the likelihood of chosen responses and reduce the likelihood of rejected responses.

10.1.2 Train on preference pairs (x, y_w, y_l) without a separate reward model.

10.2 Apply a quality gate before accepting an output:

10.2.1 Score each deliverable with the fine-tuned preference model.

10.2.2 Check instruction following, format, and safety requirements.

10.3 Handle quality failures via targeted DAG re-planning:

10.3.1 Record errors when an output fails the quality threshold.

10.3.2 Send the failed branch back to the planner for targeted revision.

Step 11: Artifact Assembly and Global State Commit

11.1 Aggregate and assemble multi-agent deliverables:

- **11.1.1** Collect verified outputs from completed DAG tasks.
- **11.1.2** Combine reports, code changes, and formatted content into one deliverable.

11.2 Evaluate the final deliverable for correctness, completeness, and timeliness:

- **11.2.1** Reject or revise the deliverable if a required dimension is weak.
- **11.2.2** Log quality results to improve future routing decisions.

11.3 Save the completed session state:

- **11.3.1** Write conversation records, DAG state, and tool receipts to `PrismDB`.
- **11.3.2** Record latency, token use, and audit information.

Step 12: Output Delivery and Reactive Workspace Presentation

12.1 Synthesize and package final deliverable artifacts:

- **12.1.1** Package reports, verified code patches, and document drafts.
- **12.1.2** Create an audit receipt with agent, timing, and tool information.

12.2 Render presentation layer via multi-pane workspace interface:

- **12.2.1** Show task states such as `Queued`, `Running`, `Success`, and `Failed`.
- **12.2.2** Render the result with formatting, syntax highlighting, and export controls.
- **12.2.3** Show approval controls for pending external actions.

12.3 Synchronize client state reactively via local database transactions:

- **12.3.1** Save verified state changes and output artifacts in `PrismDB`.
- **12.3.2** Update the interface when local database state changes.

Project Architecture, Working & System Design

Project workflow, architecture, working process, agent lifecycle, dataset pipeline, DAG scheduling, and flow diagrams

06. PrismSpace Workflow (Figure 1.0)

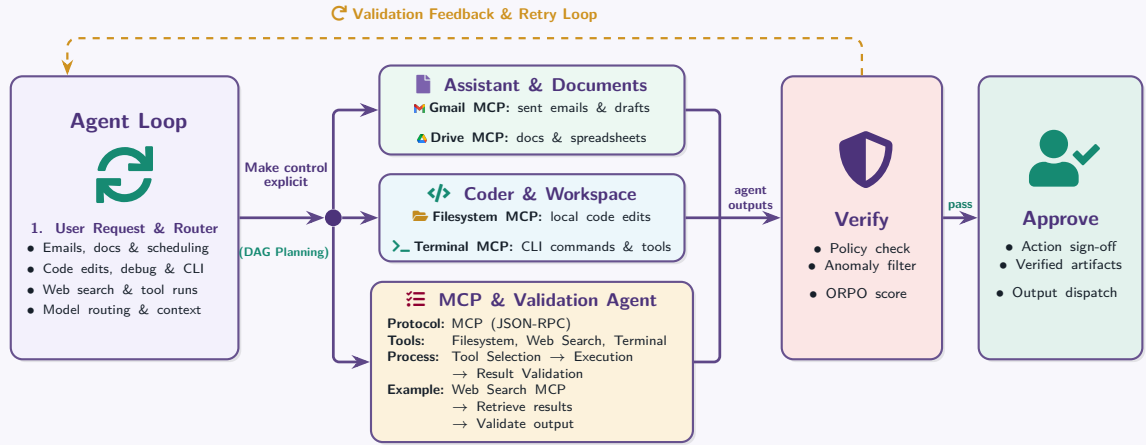


Figure 1.0: PrismSpace Workflow | **DAG** = Directed Acyclic Graph | **MCP** = Model Context Protocol | **ReAct** = Reason + Act | **ORPO** = Odds Ratio Preference Optimization | **HITL** = Human-in-the-Loop

10. Unified Dataset Preparation & Multi-Model Training Pipeline

Raw multi-source datasets flow through normalization and curation into four supervised tables, then train eleven production models.

1. Raw Datasets

- **WildGuardMix**
- **Agent trajectories**
- **RouterBench**
- **UltraFeedback**

`.parquet, .json, .jsonl`

normalize + curate

2. Supervised Tables: Fields & Training Roles

Dataset → Table	Training Fields	How Fields Are Used in Training
WildGuardMix `approval.jsonl`	`prompt, label (harm)`	Trains **ApprovalPredictor** & anomaly filter to block harmful/adversarial prompts.
Trajectories `success.jsonl`	`prompt, reward (0/1)`	Trains **SuccessPredictor** to estimate DAG task completion probability.
RouterBench `provider.jsonl`	`prompt, winner, vote_score`	Trains **ModelRouter** to select optimal provider (Groq vs OpenAI vs NIM).
UltraFeedback `reward.jsonl`	`prompt, chosen, rejected`	Trains **ORPO RewardModel** (Qwen2.5-1.5B) via preference ranking loss.

train + deploy

3. Production Models (11)

Routing & Intent Models (3)
- IntentClassifier
- AgentRouter
- ModelRouter: Groq / OpenAI / NIM

Workflow Intelligence & Performance (3)
- WorkflowClusterer
- SuccessPredictor
- LatencyPredictor

Safety, Alignment & Context (5)
- ApprovalPredictor, AnomalyDetector
- FAISS Vector Retriever
- RewardModel + Qwen2.5-1.5B ORPO

Dataset split: 80% training | 10% validation | 10% testing.
The held-out 20% is divided equally for validation and final testing.

10. Model Training Map: Algorithm Choice by Responsibility

Model-to-Algorithm Training Matrix

Model	Type	ML algorithm for training	Why choose this algorithm?
Intent Classifier	Classification	TF–IDF + XGBoost classifier	Fast and strong on sparse text features; captures nonlinear boundaries across the 12 intent classes.
Agent Router	Multilabel	Binary-relevance XGBoost	Predicts several specialist agents independently while preserving strong tabular/text-feature performance.
Provider Router	Classification	XGBoost ranking/classifier	Combines capability, quality, cost, and availability signals to choose the best provider/model.
Anomaly Detector	Anomaly	Isolation Forest	Works with few or no anomaly labels and isolates rare, unusual execution patterns efficiently.
Reward ORPO Adapter	ORPO	ORPO + Reinforcement Learning	Learns from chosen-versus-rejected preference feedback to improve response quality and alignment.
FAISS Retrieval Index	Unsupervised	Sentence embeddings + FAISS	Requires no class labels; semantic vectors support fast nearest-neighbor retrieval over workspace context.

09. Model Responsibilities & Inference Flow

Intent: The **underlying goal** of a user's request. Determines *task type* (research, coding, tool use, synthesis) for routing to agents & tools.

1. Intent Classifier	**Prompt** → **Extracted features** → **classify intent** (research, tool use, coding, everyday synthesis)	**Task type** `intent_classifier.joblib`
2. Agent Router	**Intent + task features** → predict **multiple agent labels** → form the **specialist team**	**Agent selection** `agent_router.joblib`
3. Workflow Planner	**Workflow templates** → estimate **completion probability** → choose a **dependency-safe DAG plan**	**Plan reuse + success** `workflow_templates.pkl` `workflow_success_predictor.joblib`
4. Provider Routing	**Candidate task** → choose **provider/model** → compare **predicted latency** and **token cost**	**Provider + cost balance** `model_router.joblib` `latency_predictor.joblib` `cost_predictor.joblib`
5. Safety Gate	**Planned action** → score **approval risk** + **detect anomalies** → **approve, pause, or block** execution	**Risk approval & safety** `approval_predictor.joblib` `anomaly_detector.joblib`
6. Context & Quality	**Query** → **FAISS semantic retrieval** → add **context** → generate & **ORPO-align response**	**Evidence + output quality** `faiss.index` · `faiss.vectorizer.joblib` `reward_orpo_v2/adapter_model.safetensors`

Intent Classifier: Prompt Decomposition, NLP Algorithm & Serialization

Formal Pipeline: $\mathcal{T}_{\text{user}} \xrightarrow{\text{Dense Tokenize}} \mathbf{e} \in \mathbb{R}^{384} \xrightarrow{\text{Intent Classifier}} \hat{\mathbf{y}}_{\text{intent}} \xrightarrow{\text{Semantic Decomposition}} \mathcal{G}(V, E) \xrightarrow{\text{JSON Schema}}$ JSON-RPC Task Packet

Raw Input $\mathcal{T}_{\text{user}}$

"Find the latest API changes, summarize them, and update my notes."

vector $\mathbf{e}$

Intent Classifier

- **Extracts intent:** Understands requested user actions.
- **Maps task domains:** Identifies web search, synthesis & tools.
- **Enforces safety rules:** Flags file modifications for approval.

intents $\hat{\mathbf{y}}$

Semantic Role Decomposer

- **Algorithm:** Action-Entity-Precondition SRL
- **Decomposition:** 4 Vertices (V), 3 Edges (E)
- **Rule:** Causal pipeline $v_i \rightarrow v_{i+1}$

Node v_1: Research
SearchWeb("API changes")
Browser MCP · raw_docs

raw_docs

Node v_2: Extraction
ParseDOMDiff(raw_docs)
Parser MCP · api_diff

api_diff

Node v_3: Synthesis
Summarize(api_diff)
Writer Agent · summary_md

summary

Node v_4: Approval & Write
Filesystem.Write(summary)
HITL Gate $\Longrightarrow$ Commit notes

serialize

Strongly-Typed JSON-RPC Task Packet Schema

```
{
 "task_id":  "dag_8921", "algorithm":  "kahn_topological",
 "vertices":  [
  {"id":  "v1", "agent":  "Researcher", "mcp":  "browser.search", "deps":  []},
  {"id":  "v2", "agent":  "Analyst", "mcp":  "parser.extract", "deps":  ["v1"]},
  {"id":  "v3", "agent":  "Writer", "mcp":  "llm.summarize", "deps":  ["v2"]},
  {"id":  "v4", "agent":  "Filesystem", "mcp":  "fs.write_file", "deps":  ["v3"],
"approval_gate":  true}
 ]

}
```

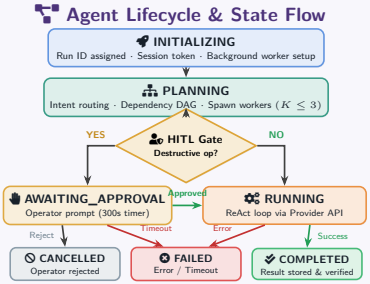
execute

Task Dispatcher

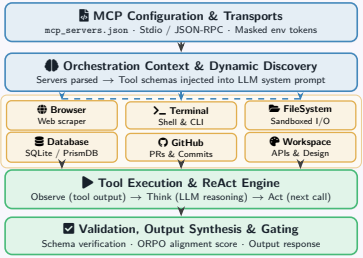
- **Runs tasks in order:** Dispatches ready tasks and passes outputs forward.
- **Enforces approval:** Pauses before writing files until user confirms.
- **Verifies quality:** Checks outputs before saving results.

07. Agent Lifecycle, MCP Integration & Runtime Flow in PrismSpace

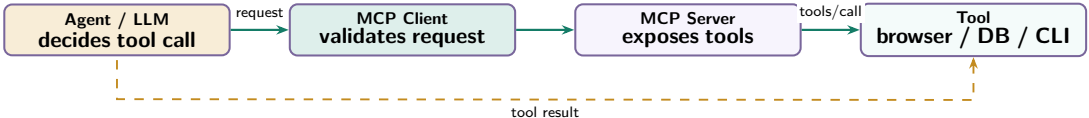
Agent Lifecycle & State Flow



MCP Tool Integration & Runtime Flow



How Agents Communicate with MCP




```
-> request:  {"jsonrpc":"2.0","id":7,"method":"tools/call","params":{"name":"search","arguments":{"q":"MCP"}}}
<- response:  {"jsonrpc":"2.0","id":7,"result":{"content":[{"type":"text","text":"3 matches"}]}}
```

11. How the System Works: From Request to Result

Stage 1: Ingestion & Features

- **Capture request:** Parse user input (text/voice/API).
- **Clean & normalize:** Remove noise, handle missing data with NLP preprocessing.
- **Extract features:** Tokenize into semantic embeddings using transformer models.

Stage 2: Intelligence & DAG

- **Classify intent:** ML classifier identifies task type from 12 categories (99.6% accuracy).
- **Choose agents & models:** Route sub-tasks to specialized LLM agents.
- **Topological sorting:** Order DAG tasks by dependencies; execute only after all prerequisites finish.

Stage 3: Execution & Storage

- **Execute:** Run up to 3 parallel tasks using Model Context Protocol (MCP) tools.
- **Track progress:** Real-time telemetry, node state transitions, and execution logs.
- **Verify & store:** ORPO quality check, persist results to PrismDB.

Building & Executing the Task Graph $\mathcal{G}(V, E)$

1. User Goal Prompt
"Analyze news+CSV, email summary"

→

2. Split into Vertices V
Modular async tasks A, B, C, D

→

3. Connect Edges E
Map output return → input args

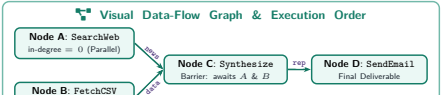
→

4. Topological Sort
Pick in-degree $= 0$; unlock dependents

Computer Science Explanation (By Concrete Example):

- **Vertices (V) = Sub-Task Functions (Async Call Units):**
  Node A: `news = await SearchWeb("Tesla news")` *[Researcher]*
  Node B: `data = await FetchCSV("metrics.csv")` *[Data Agent]*
  Node C: `rep = await Synthesize(news, data)` *[Writer Agent]*
  Node D: `await SendEmail(to="lead", body=rep)` *[Email Tool]*
- **Edges (E) = Data Pipes & Synchronization Barriers:**
  Edge (A, C) pipes return value `news`; edge (B, C) pipes `data`.
  Barrier: task C blocks until both promises A & B resolve.
- **Connection Basis:** C links to A & B because its inputs require their outputs; A and B are independent.
- **Topological Sorting (Kahn's method):** repeatedly select nodes with in-degree $= 0$, remove outgoing edges, and unlock dependents.
  This gives a valid dependency order and exposes independent work for parallel execution.
- **Execution:** A and B run in parallel; after both finish, C unlocks, then D runs.

Visual Data-Flow Graph & Execution Order



Parallel Thread 1: Worker 1 executes Node A → returns `news`
Parallel Thread 2: Worker 2 executes Node B → returns `data`
Join Barrier: Node C consumes `news` + `data` → outputs `rep`

Serial Finish: Node D consumes `rep` → sends email

Why Topological Sorting is Chosen for DAG Task Scheduling

1. Partial Order Precedence

- **Causal constraints:** Task v blocks until inputs $u \in \text{Pred}(v)$ finish.
- **No false priority:** Independent tasks (A, B) share no forced order.

2. Finds Cycles & Prevents Hangs

- **Finds circular dependencies:** Detects when tasks wait for each other in a loop $(A \to B \to A)$.
- **Stops safely:** Reports the problem instead of waiting forever for tasks to finish.

3. Parallel Concurrency Fan-out

- **Zero in-degree queue:** Nodes with in-degree $= 0$ run **concurrently**.
- **Barrier sync:** Automatically resolves dependencies into join barriers.

Comparative Analysis: Topological Sort vs. Classical Sorting Algorithms

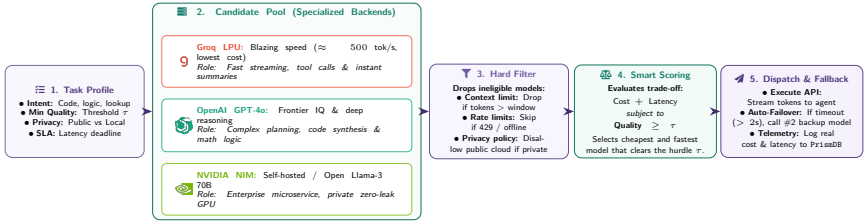
Algorithm	Data Structure	Ordering Principle	Time Compl.	Cycle Detect	Parallelism	Fit for Workflow DAGs
Topological Sort (Kahn's / DFS)	Directed Acyclic Graph $\mathcal{G}(V, E)$	**Causal Precedence** (Partial Order: $u \to v$)	$\mathcal{O}(V + E)$	Yes (In-degree)	Yes (Fan-out)	**Optimal:** Respects data pipes; runs async independent tasks in parallel.
Merge Sort (Divide & Conquer)	1D Array / Linear List	**Value Magnitude** (Total Order: $a \leq b$)	$\mathcal{O}(n \log n)$	No (N/A)	No (Serial 1D)	**Invalid:** Ignores dependency graphs; compares data values, not prerequisites.
Quick Sort (Pivot Partitioning)	1D Array / Linear List	**Value Magnitude** (Total Order: $a \leq b$)	$\mathcal{O}(n \log n)$ avg	No (N/A)	No (Serial 1D)	**Invalid:** Partitions around arbitrary value pivots; destroys execution order.
Heap Sort (Priority Queue)	Binary Tree / Array	**Key Priority** (Total Order: max/min)	$\mathcal{O}(n \log n)$	No (N/A)	No (Serial 1D)	**Invalid:** Orders by scalar priority; cannot enforce multi-parent edge barriers.
Bubble / Insertion (Linear Swap)	1D Array / Linear List	**Adjacent Swap** (Total Order: $a_i > a_{i+1}$)	$\mathcal{O}(n^2)$	No (N/A)	No (Strict Serial)	**Invalid:** Quadratic pairwise swaps cannot represent topological DAG flow.

Core Architectural Takeaway: Classical sorting algorithms ($\mathcal{O}(n \log n)$) arrange scalar values along a 1D axis ($a_1 \leq a_2 \ldots$). In contrast, multi-agent execution graphs require scheduling a **partially ordered set** with zero, one, or multiple asynchronous parent promises. **Topological Sorting** is the only algorithm with $\mathcal{O}(V + E)$ complexity that guarantees causal precedence, detects deadlocks, and unlocks parallel execution.

How Provider Routing Works: Task-to-Model Selection

Core Routing Mission: Instead of sending every query to a slow, expensive frontier model, PrismSpace dynamically routes each agent sub-task to the **fastest, most cost-effective provider** that clears the task's required **quality threshold (τ)**.

End-to-End Provider Routing Pipeline



How the Router Decides (Multi-Objective Optimization)

Objective: $\min_m \left(\lambda_c \cdot \widehat{\text{Cost}}_m + \lambda_l \cdot \widehat{\text{Latency}}_m \right)$ s.t. $\widehat{\text{Quality}}_m \geq \tau$

- **Speed-Critical Nodes ($\lambda_l \uparrow$):** Fast tool calls prioritize low latency $\rightarrow$ routes to **Groq LPU** (< 300ms).
- **High-Complexity Tasks ($\tau \uparrow$):** Architectural synthesis prioritizes IQ $\rightarrow$ routes to **OpenAI GPT-4o**.
- **Confidential Workflows:** Data protection requirements lock execution to local **NVIDIA NIM** microservices.

Runtime Safeguards & Zero-Downtime Fallback

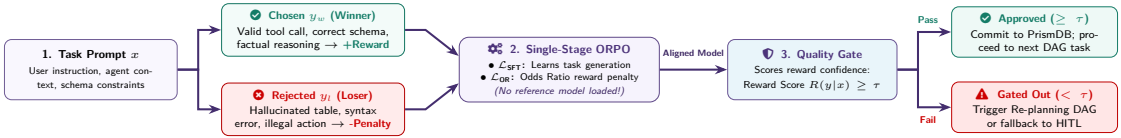
- **Pre-Flight Verification:** Validates endpoint health, remaining quota, and network jitter before dispatching tokens.
- **Dynamic Circuit Breaker:** If the primary provider times out (> 2000ms) or hits HTTP 429/503 rate limits, execution instantly cascades to the #2 ranked model (< 150ms failover).
- **Closed-Loop Feedback:** Logs observed token speed, dollar cost, and quality scores to continuously calibrate future routing models.

How ORPO Works: Preference Alignment and Quality Gating

What is ORPO & Why is it Used Here as a Reward System?

- **What is ORPO? Odds Ratio Preference Optimization (ORPO)** is an integrated **single-stage reward and alignment algorithm**. Traditional RLHF (PPO) requires training a separate reward model and keeping a frozen reference model in GPU memory (heavy, slow, and unstable).
- **Why in PrismSpace?** In multi-agent swarms, LLMs can hallucinate invalid tools or malformed code. ORPO acts as an **internal reward filter**: it mathematically rewards compliant trajectories (y_w) and penalizes bad responses (y_l) during fine-tuning, eliminating the need for an external reward model.

End-to-End Pipeline: From Preference Pairs to Runtime Quality Gating



ORPO vs. Traditional RLHF

- **The RLHF (Reinforcement Learning from Human Feedback) Problem:** Traditional alignment (PPO/DPO) requires running 3 separate models simultaneously: Policy Model, Reference Model, and Reward Critic. This demands > 48 GB VRAM and suffers from unstable policy gradients.
- **The ORPO Breakthrough:** It mathematically merges preference alignment directly into supervised fine-tuning in a **single stage**.
- **Concrete Benefit: 50% GPU memory savings**, faster convergence, and no policy drift without loading any secondary reference model!

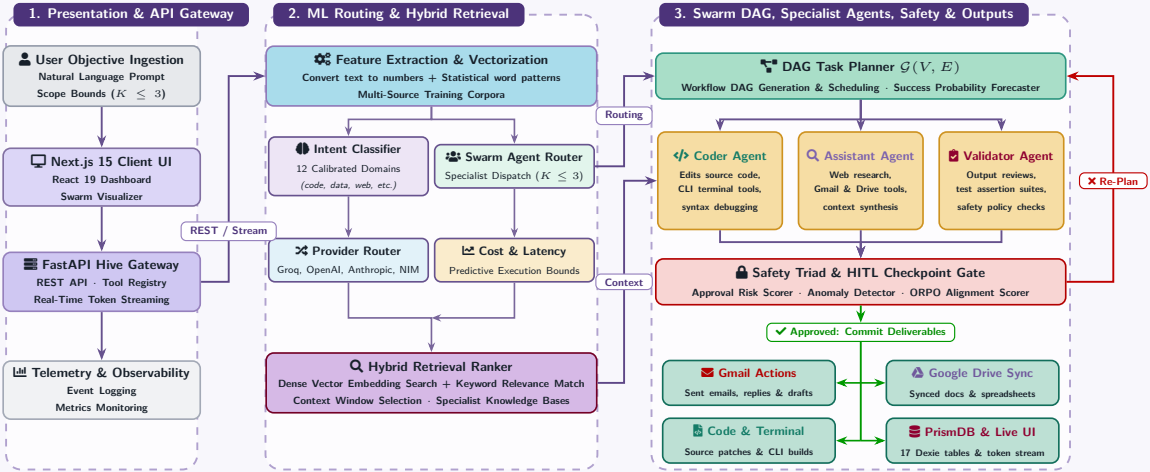
The Mathematical Reward Intuition

$$\mathcal{L}_{\text{ORPO}} = \mathcal{L}_{\text{SFT}} + \lambda_{\text{OR}}\, \mathcal{L}_{\text{OR}}$$

$$\mathcal{L}_{\text{OR}} = -\log \sigma\left(\log \frac{\text{odds}(y_w|x)}{\text{odds}(y_l|x)}\right), \quad \text{odds}(y|x) = \frac{P(y|x)}{1-P(y|x)}$$

- **Implicit Reward Gradient:** $\text{odds}(y|x)$ measures the probability of generating a token vs not generating it. When y_l is chosen, loss explodes, forcing logits of $y_w \gg y_l$.
- **Runtime Impact:** The model naturally generates schema-compliant tool calls and suppresses hallucinations before writing to PrismDB.

12. PrismSpace System Architecture & Multi-Model Swarm Workflow (Figure 1.2)



13. Flowchart: Multi-Agent Execution Flow (Figure 1.3)

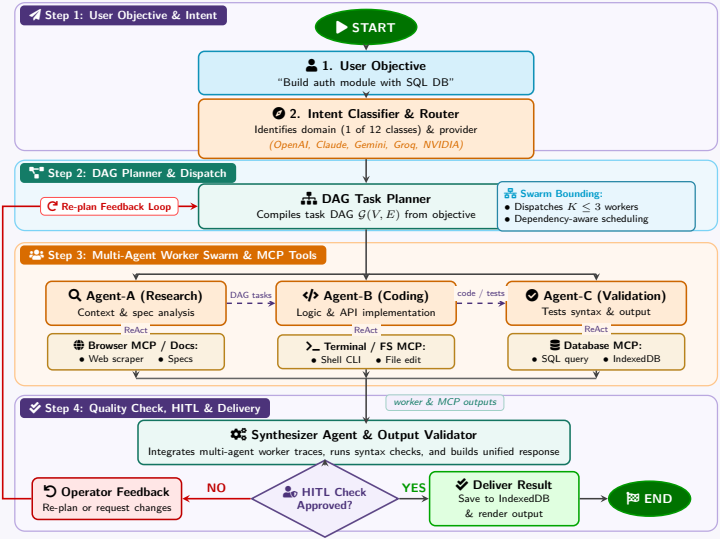
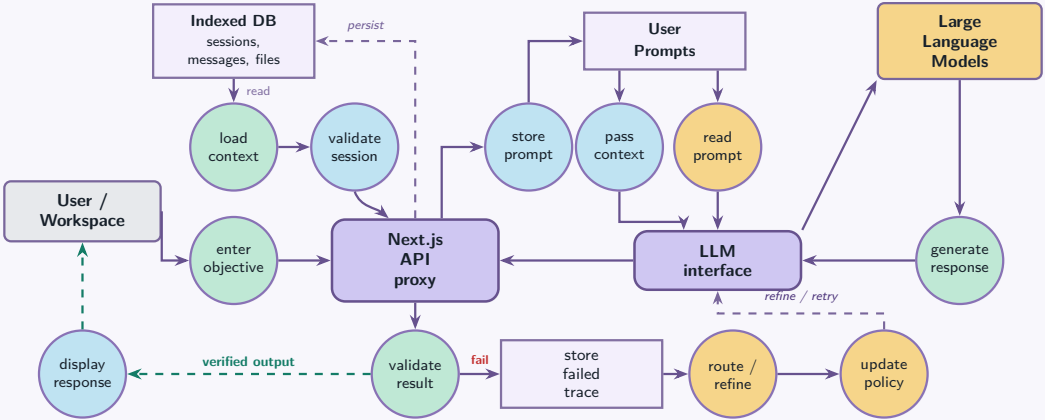


Figure 1.3: Simplified Multi-Agent Communication Flow in PrismSpace

07. PrismSpace Data Flow Diagram



Database Design

PrismDB schema, relationships, and entity–relationship structure

14. Database Schema

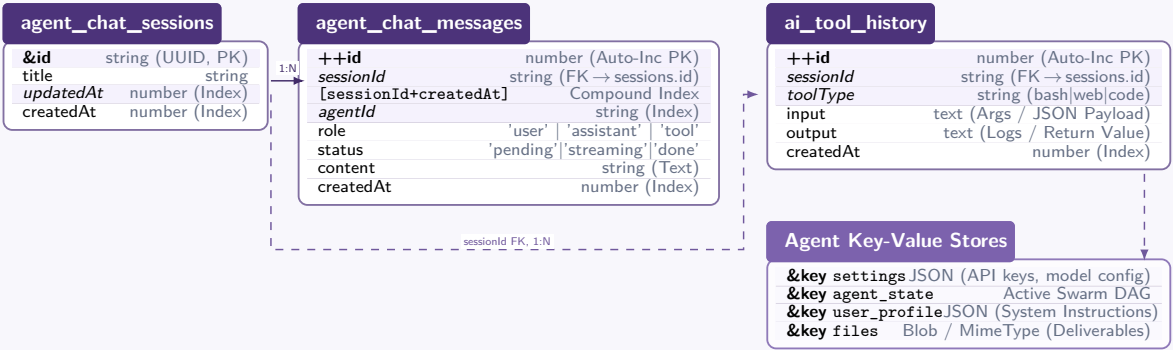
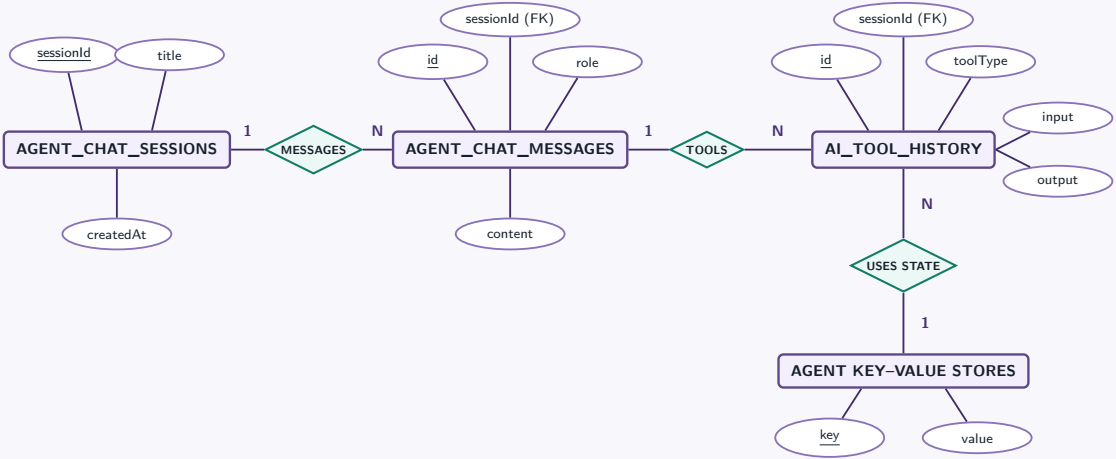


Figure 1.4: PrismDB - AI-Agent Browser IndexedDB Schema (Dexie.js v3) | **PK**: Primary Key **FK**: Foreign Key **Index**: B-tree index
`sessions.id` (Shared FK) → `messages.sessionId` & `tool_history.sessionId` | `useLiveQuery()` Reactive UI | ACID Dexie.js

15. Database ER Diagram



Project Requirements

Hardware and software needed to build and run PrismSpace

16. Hardware Requirements

- **CPU:** Multi-core processor for the Next.js client, FastAPI service, and local development tools.
- **GPU:** NVIDIA RTX 3050 or above for embedding generation, vector indexing, and model-training workloads.
- **Memory:** Minimum 16 GB RAM for running the frontend, backend, databases, and supporting services concurrently.
- **Storage:** SSD with sufficient capacity for the codebase, PrismDB data, logs, model checkpoints, and FAISS indexes.
- **Network:** Reliable internet connection for provider APIs, MCP integrations, and model downloads.
- **Security:** Secure environment-variable or secret-manager support for provider and tool credentials.

17. Software Requirements

- **Next.js 14 / React / TypeScript:** browser-based workspace for objectives, chat history, agent status, and approvals.
- **FastAPI / Python:** asynchronous backend framework for agent orchestration, REST endpoints, and runtime events.
- **Dexie.js / IndexedDB:** local ACID storage for sessions, messages, files, settings, and tool history.
- **Joblib / scikit-learn:** Python libraries for data preparation, utility functions, serialization, and general machine-learning support.
- **MCP JSON-RPC:** MCP (Model Context Protocol) uses JSON-RPC (JavaScript Object Notation–Remote Procedure Call) messages to standardize tool names, inputs, outputs, and errors for browser, terminal, filesystem, GitHub, database, and workspace tools.
- **Provider SDKs/APIs: NVIDIA NIM** provides optimized model-serving endpoints; the **Groq SDK** supports low-latency inference; and the **Anthropic Claude SDK/API**, OpenAI-compatible APIs, and similar provider clients connect the orchestrator to different language models.

Project Implementation

Frontend, backend, and training components

18. Frontend Implementation

- **Next.js 14:** React-based framework for building the browser application and managing frontend pages.
- **React:** component library for creating reusable user-interface elements.
- **TypeScript:** typed JavaScript used for safer and more maintainable frontend development.
- **Tailwind CSS:** utility-first styling framework for responsive layouts and consistent visual design.
- **Dexie.js, IndexedDB, and dexie-react-hooks:** browser storage and reactive data-access libraries for local application data.
- **Lucide React:** icon library used for interface symbols and controls.
- **Framer Motion and GSAP (GreenSock Animation Platform):** frontend animation libraries for smooth interface transitions.

19. Backend Implementation

- **FastAPI:** Python framework for building asynchronous backend services and application endpoints.
- **Python:** primary backend language for API services, integrations, data handling, and runtime coordination.
- **Uvicorn:** lightweight **ASGI (Asynchronous Server Gateway Interface)** server used to run FastAPI. ASGI connects asynchronous Python applications to web servers and supports non-blocking concurrent requests.
- **REST (Representational State Transfer):** API style for communicating with the frontend through structured HTTP requests and responses.
- **Pydantic:** schema validation and typed data models for request, response, configuration, and tool payloads.
- **MCP (Model Context Protocol) and JSON-RPC (JavaScript Object Notation–Remote Procedure Call):** standards for structured communication between the backend and external tools.
- **Provider SDKs/APIs:** integration layer for connecting the backend to external language-model and model-serving services.

20. Model Training Modules

- **Dataset preparation:** collects, cleans, and combines task traces and preference data using **Pandas**, **NumPy**, and **PyArrow**.
- **Intent classification:** categorizes user objectives using **scikit-learn** text features, classifiers, and **XGBoost**.
- **Agent routing:** selects specialist agents using shared **scikit-learn** pipelines and **XGBoost** classification support.
- **Model and provider routing:** compares task features and provider signals using **Pandas**, **NumPy**, **scikit-learn**, and **XGBoost**.
- **Performance prediction:** estimates latency and cost with **XGBoost** regression, supported by **scikit-learn** preprocessing.
- **Safety prediction:** estimates approval needs and detects unusual patterns with **scikit-learn**, **XGBoost**, and **Isolation Forest** utilities.
- **Retrieval:** uses **FAISS (Facebook AI Similarity Search)**, **NumPy**, and optional **sentence-transformers** embeddings for context search.
- **Preference alignment:** uses **PyTorch**, **Hugging Face Transformers**, **TRL (Transformer Reinforcement Learning)**, **PEFT**, and **Datasets** for ORPO training.
- **Runtime artifacts:** uses **Joblib**, **Safetensors**, and **Hugging Face Accelerate** to save, validate, and load trained components.

Artifact: a saved output from training, such as a trained model, vector index, or configuration file, that can be reused later. **FAISS (Facebook AI Similarity Search):** a library for quickly finding similar vectors. **PEFT (Parameter-Efficient Fine-Tuning):** adapts a model by training only a small set of additional parameters.

System Design

Modules, responsibilities, and system-level organization

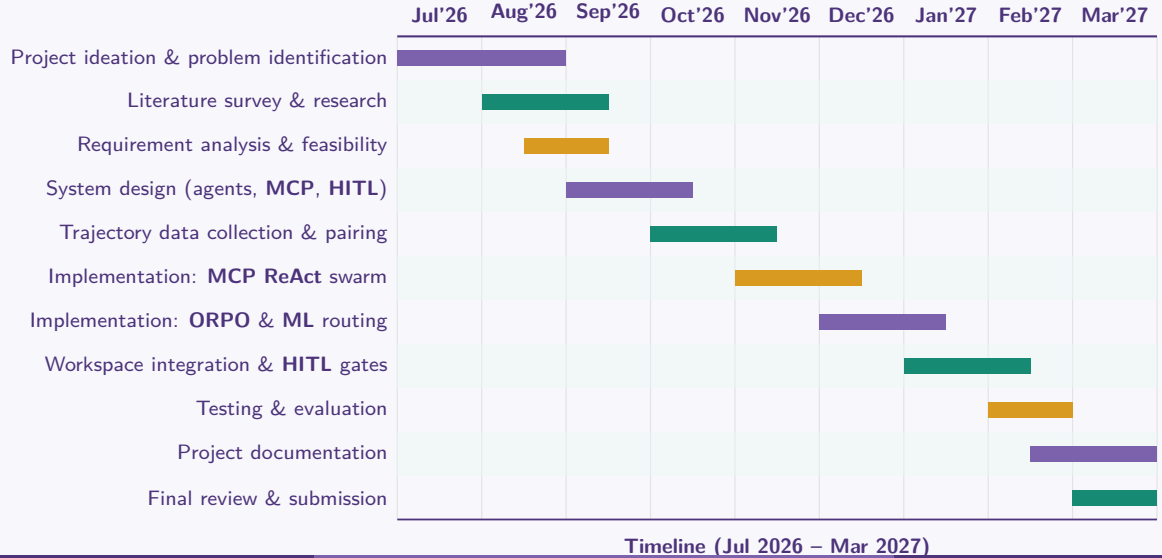
21. System Modules and Responsibilities

- **User and Session Module:** captures objectives, settings, conversation history, uploaded files, and session state.
- **Context and Routing Module:** classifies intent and selects the most suitable provider, model, and specialist agents.
- **DAG (Directed Acyclic Graph) Planning Module:** decomposes objectives into dependency-aware subtasks and schedules parallel work.
- **Agent Execution Module:** runs bounded **ReAct (Reasoning and Acting)** loops and invokes approved tools.
- **MCP (Model Context Protocol) Integration Module:** exposes filesystem, terminal, database, memory, and external-service tools through **JSON-RPC (JavaScript Object Notation–Remote Procedure Call)**.
- **Verification and Safety Module:** applies schema checks, **ORPO (Odds Ratio Preference Optimization)** preference gating, anomaly detection, and **HITL (Human-in-the-Loop)** approvals.
- **Persistence Module:** stores sessions, messages, artifacts, and errors in the local **PrismDB (browser IndexedDB database)**.

22. Conclusion

- **PrismSpace** converts a user's natural-language goal into smaller, organized tasks.
- Independent tasks can be executed by multiple agents in parallel.
- Standard tools are connected through the **Model Context Protocol (MCP)**.
- Human approval and local browser storage improve safety, privacy, and control.
- The platform provides a flexible foundation for future agent-based applications.

23. Gantt Chart



References

[1] **Cornalba F, Disselkamp C, Scassola D, Helf C (2024)** Multi-objective reward generalization: improving performance of Deep Reinforcement Learning for applications in single-asset trading. *Neural Computing and Applications*. DOI: 10.1007/s00521-023-09033-7

[2] **Schneckenreither M, Moser G (2025)** Average reward adjusted discounted reinforcement learning. *Neural Computing and Applications*. DOI: 10.1007/s00521-024-10620-5

[3] **Aina K, Ha S (2026)** Deep reinforcement learning for multi-agent coordination. *Artificial Life and Robotics*. DOI: 10.1007/s10015-025-01089-z

[4] **Rother D, Pajarinen J, Peters J, Weisswange TH (2025)** Open-ended coordination for multi-agent systems using modular open policies. *Autonomous Agents and Multi-Agent Systems*. DOI: 10.1007/s10458-025-09723-7

[5] **Liu D, Ren F, Yan J, Su G, Gu W, Kato S (2026)** Improving scalability of multi-agent deep reinforcement learning with suboptimal human knowledge. *Autonomous Agents and Multi-Agent Systems*. DOI: 10.1007/s10458-025-09729-1

[6] **Marzari L, Liu C, Donti PL, Marchesini E (2026)** ε-retraining reinforcement learning algorithms. *Autonomous Agents and Multi-Agent Systems*. DOI: 10.1007/s10458-026-09748-6

[7] **Wong A, Bäck T, Kononova AV, Plaat A (2023)** Deep multiagent reinforcement learning: challenges and directions. *Artificial Intelligence Review*. DOI: 10.1007/s10462-022-10299-x

[8] **Niu B, Luo B, Cui Y, Xu X, Zhao Y, Feng Y (2025)** HIAT: human-in-the-loop reinforcement learning with auxiliary task. *Artificial Intelligence Review*. DOI: 10.1007/s10462-025-11260-4

References

[9] **Lee J (2025)** InstructPatentGPT: training patent language models to follow instructions with human feedback. *Artificial Intelligence and Law*. DOI: 10.1007/s10506-024-09401-1

[10] **Zhuang S, Qiao Z, Zuccon G (2023)** Reinforcement online learning to rank with unbiased reward shaping. *Information Retrieval Journal*. DOI: 10.1007/s10791-022-09413-y

[11] **Ghazal TM, Awouda AE, Hasan MK, Rahman AHA, Islam S, Saeed RA, Elshafie H, Iqbal A, Hussein AH (2025)** Knowledge Learning for Securing Workflow Scheduling Algorithm in Mobile Edge Computing. *Mobile Networks and Applications*. DOI: 10.1007/s11036-025-02465-6

[12] **Barredo P, Puente J (2025)** Energy-aware cooperative multi-fitness evolutionary algorithm for workflow scheduling in cloud computing. *Natural Computing*. DOI: 10.1007/s11047-025-10023-y

[13] **Heirati A, Fartash M, Khalily-Dermany M (2025)** Optimized Task Scheduling in Fog-Cloud Computing Using Hybrid Deep Learning and Metaheuristic Algorithms. *Neural Processing Letters*. DOI: 10.1007/s11063-025-11819-w

[14] **Abdi S, Ashjaei M, Mubeen S (2024)** Cost-aware workflow offloading in edge-cloud computing using a genetic algorithm. *The Journal of Supercomputing*. DOI: 10.1007/s11227-024-06341-0

[15] **Xia Y, Zhu J, Zhu L (2023)** Dynamic role discovery and assignment in multi-agent task decomposition. *Complex & Intelligent Systems*. DOI: 10.1007/s40747-023-01071-x

[16] **Li J, Xing L, Zhong W, Cai Z, Hou F (2023)** Decision variable contribution based adaptive mechanism for evolutionary multi-objective cloud workflow scheduling. *Complex & Intelligent Systems*. DOI: 10.1007/s40747-023-01137-w

References

[17] **Yu X, Lin Y, Wang X, Han S, Lv K (2024)** GHQ: grouped hybrid Q-learning for cooperative heterogeneous multi-agent reinforcement learning. *Complex & Intelligent Systems*. DOI: 10.1007/s40747-024-01415-1

[18] **Xu C, Wang J, Zhu X, Yue Y, Zhou W, Liang Z, Wojtczak D (2024)** Decentralized multi-agent cooperation via adaptive partner modeling. *Complex & Intelligent Systems*. DOI: 10.1007/s40747-024-01421-3

[19] **Vyhmeister E, Castañé GG, Östberg PO, Thevenin S (2023)** A responsible AI framework: pipeline contextualisation. *AI and Ethics*. DOI: 10.1007/s43681-022-00154-8

[20] **Cavalcante Siebert L, Lupetti ML, Aizenberg E, Beckers N, Zgonnikov A, Veluwenkamp H, Abbink D, Giaccardi E, Houben GJ, Jonker CM, van den Hoven J, Forster D, Lagendijk RL (2023)** Meaningful human control: actionable properties for AI system development. *AI and Ethics*. DOI: 10.1007/s43681-022-00167-3

[21] **Yue Q, Tian L, Feng X, Yuan Z, Wei S, Xia J, Chen B, Song X, Yao Y (2026)** CPN-HRL: a hierarchical deep reinforcement learning approach for priority-aware task scheduling in CPN enabled by cloud-edge-end environments. *Journal of King Saud University Computer and Information Sciences*. DOI: 10.1007/s44443-026-00690-x

[22] **Skirzyński J, Jain Y, Lieder F (2023)** Automatic discovery and description of human planning strategies. *Behavior Research Methods*. DOI: 10.3758/s13428-023-02062-z

[23] **Zhang S, Zhao Z, Liu C, Qin S (2023)** Data-intensive workflow scheduling strategy based on deep reinforcement learning in multi-clouds. *Journal of Cloud Computing*. DOI: 10.1186/s13677-023-00504-9

[24] **Wang Z, Chen S, Bai L, Gao J, Tao J, Bond RR, Mulvenna MD (2023)** Reinforcement learning based task scheduling for environmentally sustainable federated cloud computing. *Journal of Cloud Computing*. DOI: 10.1186/s13677-023-00553-0

[25] **Subashree S, Rajakumaran M, Pushpa G, Senthilkumar S (2025)** Hybrid dolphin swarm sparrow search optimization based multi-objective cloud workflow scheduling. *Journal of Cloud Computing*. DOI: 10.1186/s13677-025-00828-8

Thank You!

Nobin Sijo | Pranav P | Prasanth P | Sankar Narayan S

Department of Computer Science & Engineering
College of Engineering, Karunagappally